

# Hardware Transactional Memory without the Hardware

Zachary Bischof  
Northwestern University  
Evanston, IL, USA  
zbischof@u.northwestern.edu

Marcel Flores  
Northwestern University  
Evanston, IL, USA  
marcel-flores@u.northwestern.edu

Maciej Swiech  
Northwestern University  
Evanston, IL, USA  
maciejswiech2007@u.northwestern.edu \*

## ABSTRACT

Transactional memory has long been considered as a useful alternative to locks for implementing parallel algorithms, due to its simpler programming interface. In particular, it offers an optimistic approach to parallelization, and attempts to perform an operation assuming no contention will occur. Despite the potential gains of hardware transactional memory, it is not yet widely supported in hardware. This will change soon with Intel's upcoming Haswell architecture, which provides support for certain transactions. However, there is not an easy way to measure how the added support for these instructions will help improve performance.

In this paper, we leverage Palacios, a virtual machine monitor, to add preliminary support for Intel's publicly available transactional memory specification. By switching between direct execution on hardware and simulated support of hardware transactional memory, we are able to move toward testing the impact of these instructions on application performance, providing developers with a unique chance to estimate how hardware support could affect their application's performance. The focus of this work is to explain how we were able to add support to Palacios to test these features.

## 1. INTRODUCTION

As the level of hardware parallelism has increased, so has the complexity of concurrency control mechanisms to which programmers must be exposed. In order to fully utilize the hardware, one must juggle a growing number of locks, condition variables, and all the overhead that comes with them.

One proposed solution to this problem has been the concept of transactional memory [3], which proposes a sys-

tem which hides the management of concurrency from programmers and instead presents them with an interface which guarantees that either a segment of code will be run without interference from other processes, or will be abandoned and the program allowed to determine its next course of action.

We present a preliminary implementation of the transactional memory system described in the Intel extension for the upcoming Haswell processor [4]. In order to do so, we make use of the Palacios virtual machine monitor. In particular, our system is transparent; from the perspective of the guest operating system, it appears that the transaction occurred in hardware as if it had been run on a processor supporting the extension.

In addition to providing an example implementation of transactional memory, this work stands as an indicator of the power and flexibility of virtual machine monitors as a platform for implementing hardware design concepts without the difficulty of having actual hardware. While the implementation is by no means trivial, its cost is significantly less than that necessary to design and build hardware.

## 2. BACKGROUND

Transactional memory has long been proposed as a solution to the programmatic complexity that arises from attempting to use mutual exclusion locks in real world implementations of parallel algorithms. We will now consider a basic overview of the general idea of transactional memory. We then go on to describe the implementation that Intel intends to include in the extension for the upcoming Haswell processors.

### 2.1 Transactions

We define a *transaction* to be a sequence of memory operations that are performed by a processes such that, to all other processes, the operations seem to have been performed atomically. Furthermore, the operations must appear to the process performing the transaction as having happened one after another, with no interference from other processes. When a transaction is finished, it either commits the results of its operations to main memory, or it aborts the transaction and throws away its changes, depending on the results of a memory validation. The idea of memory transactions of this form arise from the notion of database transactions, except without the requirement that transac-

---

\*Authors are listed in alphabetical order.

	External Read	External Write
Read	Continue	Abort
Write	Abort	Abort

**Table 1: The results of a transaction given the combination of a transaction operation and an external operation on the same memory address.**

tions are durable to crashes.

In order to protect the appearance that the operations have executed serially without interference, the validation must detect if other processes have written to memory that has been used by the transaction. Likewise, in order to preserve the atomicity, the transaction must not allow other processes to interfere with the commit. Therefore, the validation must detect if other processes have written to memory that is read by the transaction (and thereby interfering with the appearance of serial execution), or other processes read memory written by the transaction (potentially interfering with atomicity). Upon an abort, a transaction may decide to attempt to run again, or revert to traditional mutex locking mechanisms. [3]

In addition to memory conflicts with other user processes, transactions may also be aborted in the event of interrupts, context switches, or other forced changes in the control flow, as they would violate the requirement that the transaction happen entirely serially.

Transactions of this form can be thought of as an optimistic approach to concurrency, contrasting with the more pessimistic approach of locks. Rather than explicitly preventing other processes from accessing memory, the transaction assumes that it will probably be able to complete unencumbered. If a conflict does occur, the transaction is able to deal with it accordingly. Locks, on the other hand, assume that conflicts are likely to occur, and directly block other processes from accessing the relevant portions of memory. Transactions therefor stand to offer a performance gain, as forced serialization can be avoided in many cases.

In addition, transactional memory stands to decrease the complexity of multi-process applications. Programmers would no longer be weighed down by keeping track of which process is holding which locks, and no longer needs to worry about deadlocks, starvation, and other side effects that can result from incorrect use of locks.

## 2.2 Intel Implementation

The starting with the Haswell, Intel has included the Intel Transactional Synchronization Extension (TSX), which is their implementation of transactional memory. The extension includes two interfaces: the first is Hardware Lock Elision, a system designed to work on legacy processors, and of no further interest to us. The second is Restricted Transactional Memory (RTM). This includes instructions for beginning a transactional region (and specifying the code to run in the event of a failure), aborting a transaction, and ending a transactional region. [4, 6]

The Intel implementation generally follows the design of [3], with a few exceptions. Most notably, each transaction is only attempted a single time: in the event of a failure or abort, the code immediately resorts to the abort handler. Additionally, the transactions are checked at a granularity of cache lines, rather than actual memory address. In order to validate each transaction at the time of completion, this implementation stores a read set, the locations of memory reads from the transaction, and the write set. An abort is triggered if another processor reads from an address that appears in the write set or writes to an address in either set.

Since this is a hardware implementation, transactions may also abort if the read and write sets exceed a certain, hardware dependent, capacity. The transaction will also be aborted for any number of instructions that may interfere with control flow: such as the writing of control registers, any variety of interrupt, processor state changes, TLB control, and so on. The specification also suggests that aborts may occur in the case of self modifying code, or other unpredictable scenarios.

The Intel implementation also offers support for nested paging, however we save discussion of this topic for later work. (Maybe just don't mention this).

The Palacios virtual machine manager offers us a unique opportunity to implement transactional memory which emulates both the original design of transactional memory, as well as a number of the subtleties observed in Intel's implementation.

## 3. RELATED WORK

The idea for transactional as a lock-free programming model first appeared in [3]. Since this original proposal, much of the discussion surrounding transactional memory has centered around questions of efficiency and design.


In [7] Shavit and Touitou propose a Software Transactional Memory, a software method of supporting transactions. Later iterations have sought to increase the capabilities of transactions, eliminating dependencies on certain hardware limitations and enabling large and unbounded transactions [1, 2, 5]. While there is value in many of these designs, we primarily focus on the original proposal of transactional memory and the planned Intel implementation.

## 4. PALACIOS IMPLEMENTATION

In order to implement transactional memory in Palacios, we attempt to follow Intel's Haswell specifications, providing us with a consistent interface with which to use transactions, as well as making future comparisons to direct hardware implementations easier. The implementation relies on storing all information that would normally be kept by the processor in the data structures available to Palacios. As a result, a great deal of the complexity of our implementation arises from our need to exit the guest and perform actions in the VMM. We now consider the primary components of our implementation.

### 4.1 State Machine

Our system is controlled by a state machine which is controlled by two variables. The first is the transaction mode,



**Figure 1: The state machine which controls how our implementation records transactions. #PF indicates that a page fault has occurred, FA is the faulting address, and write indicates the error code.**

which indicates that there is either no transaction running (OFF), a transaction running in the current core (ON), or a transaction running on a different core (OTHER\_CORE). Additionally, there is a state variable, which indicates which handler should be executed on the next exit to Palacios. These states are described in Table 2, and their exact use is described in the following sections.

#### 4.1.1 Beginning a Transaction

The transaction begins when the VM encounters an XBEGIN instruction. Since the VM is running on hardware which does not yet support this instruction, this causes the VM to exit. On an AMD machine, the SVM handler then indicates that the undefined exception was indeed the cause for the exit. An undefined exception handler is then called, which performs a limited decoding to see if the exit was caused by one of the three Haswell transactional memory instructions.

If the instruction is an XBEGIN, the VMM begins the process of launching the transaction. First, the argument to XBEGIN is stored, as this indicates the location of the fail-safe code. A vmcall is then registered with the VMM (we further describe this later). Next the various read and write set data structures are initialised, the transaction state is set to “IFETCH”, and the RIP is advanced to the next instruction.

Next, we invalidate the current shadow page table. This will cause the VM to exit to Palacios in the event of any memory access, allowing us to record it. Since we are only eliminating the shadow page table, we are not actually tampering with the state of the guest, which allows this manipulation to go undetected by the guest.

#### 4.1.2 Handling an Ifetch

Since we have eliminated the shadow page tables, as soon as the VM resumes executing, a pagefault is caused when attempting to read the instruction. Since the transaction mode has been set to ON, we are able to catch this pagefault, and in particular know both the faulting address and the associated error code. We confirm that this fault is caused by an ifetch through both the error code and by explicitly checking that the RIP and faulting address match. We advance the transaction state to “EXEC”, as we are now

executing a transaction.

In addition, we decode the current instruction in order to determine the length of its operands. We then replace the next 10 bytes of instructions with a VMMCALL. This will allow us to exit back into the VM after the current instruction has run and perform any updates as necessary. This is particularly important in the event that an instruction has no memory accesses, so that we can be sure we are handling each instruction individually. After we do so, we return control to the VM.

#### 4.1.3 Handling a Memory Access

If we observe a page fault caused in user space while our machine is in the EXEC state, we know that this is a memory access in our transaction. In the event that the error indicates the fault was caused by a read, we first check our write-set, to see if the address has previously been changed by our transaction. If not, we add the page to the read set, then allow the PTE handler to create the entry normally. However, prior to returning to the VM, we mark the page as read only, so that attempting to write to it will cause an additional fault.

If, on the other hand, the address was in our write set, we feed the PTE handler a special staging page where we store writes. The read then continues on, using and memory updates performed earlier in the transaction.

If the initial page fault is caused by a write (and in particular, possibly caused by an attempt to write on a page that is marked read only), we add the page to our write set. We then pass a special staging page to the PTE handler so that the write is performed on our staging page. This allows the writes to appear serially to the process running the transaction, but be invisible to all other processes.

Notably, commands which perform a write and a read on the same memory location will leave the read undetected, as it will not cause a page fault. However, this will not cause any problems, as a write is strictly worse than a read. In other words, any situation that would cause an abort with an address in the read set will also cause an abort for an address in the write set.

#### 4.1.4 Finishing an Instruction

After the memory operations of a given instruction have completed, the processor will encounter the VMMCall injected during the ifetch. This will call the registered handler. This handler will first restore the code that was overwritten in order to inject the VMMCall. It will then return the state of the transaction to “IFETCH”. Finally, the shadow page table is again invalidated, to guarantee that the next ifetch, and following memory access will cause page faults.

#### 4.1.5 Committing An Instruction

If an XEND instruction is encountered, the ifetch step happens exactly as above. However, when the VM attempts to execute the instruction, it will exit the VM and run the undefined exception handler, as with the XBEGIN above. When an XEND is seen, transaction mode is switched to OFF, the memory in the staging page is copied to main memory, and control is returned to the VM.

State	Description
IFETCH	The system is currently expecting a page fault caused by an ifetch of the next instruction.
EXEC	The system is expecting the next exit to be the result of a page fault for reading or writing data.
ABORT	The system has encountered an abort condition, and will exit the next time it encounters the exiting hypercall.

**Table 2: The states used by our implementation to determine how to handle an exit to Palacios.**

The next instruction will be the injected VMMCall, which will then restore the replaced code, unregister the hypercall, and return control to the VM without invalidating the page tables.

#### 4.1.6 Aborting a Transaction

As in the hardware implementation, there are numerous reasons why a transaction may need to be aborted. In addition to the preservation of the serial execution and atomic commit properties, anything which causes a context switch or general change to control flow should cause an abort. Additionally the transaction may abort itself, by calling XABORT. In all of these situations, the state is simply switched to “ABORT”. When the VMMCall runs, it will see that the transaction is slated for abortion, clean up the code injection from the ifetch, and call an abort handler. This handler reasonably disposes of the transaction data structures and sets the RIP of the VM to the failsafe code, thereby discarding any changes that the transaction made.

#### 4.1.7 External Aborts

TODO: This I guess. It is also possible that external processes will trigger an abort by interacting with the read write sets of the transacting process. Therefore, while a transaction is running in one processes, all other cores have their transaction mode set to OTHER. While this is on, the reads and writes of these processes are stored using the same techniques as for the transacting process, except that all reads and writes are done to main memory instead of a staging page. If an intersection is ever detected between the transactions write set and the read and write sets of other cores, or the read set of the transaction and others write sets, the transaction is aborted, using the same strategy used for other aborts.

## 4.2 Current Assumptions

The current implementation in palacios is subject to a number of assumptions that are currently the result of time limitations, and not intended to be a part of the final transaction model.

First, due to simplicity, the undefined exception handler only runs on AMD processors which provide the SVM extension. Additionally, the injected hypercalls are currently using the AMD model. Both of these are likely to be trivial to implement via their Intel counterparts.

In order to perform the hypercall code injection, the decoder is used to determine the length of the current instruction, so that the injected code can be placed without damaging necessary operands. However, the built in decoder currently only works with a limited set of instructions, preventing arbitrary instructions from being run in a transaction, limiting them to a small set of test instructions. While this issue can

likely be solved by switching to the QUIX86 decoder [8], as it can handle a much larger instruction set, bugs in Palacios prevented us from doing so at this time.

The current implementation also assumes that the set of instructions in a transaction occur on the same page, and do not run within 10 bytes of the end of the page. While this should be fixable by checking the alignment of the instruction address, we have left this assumption in the interest of time.

The current staging page used to store the writes done by a transaction also makes two assumptions. First, we require that all memory written by a transaction fit on a single page. While this is not unreasonable, given that transactions are intended, in general, to be smaller, it does create an artificial limitation. Second, since the handler which deals with the case of a transaction write simply passes the usual PTE handler the same staging page no matter what, it assumes that the offset will not conflict with any previous or future writes. We see that this could happen if two writes which normally happen on separate pages, but the same offset, are performed in the same transaction. Again, while this scenario is unlikely for small transactions we are creating a restriction.

During testing, we observed that after every reentry to the VM, a number of interrupts occur, and control jumps to the kernel. While we have not yet determined the cause of these interrupts, we suspect they are the result of a timer, or other piece of hardware. Since these may interfere with our transaction, we should abort when this occurs. However, since our test transactions are simple, we assume that they will not interfere with a transaction and simply ignore the interrupts. In particular, this means that we ignore a wave of page faults on kernel address upon each return to the VM. Since the error code of the fault indicates the faults were raised by the kernel, we can easily filter them out.

The transaction mode indicator currently assumes that only a single transaction can run at a time. This means that if one process is in transaction mode, a second core encountering an XBEGIN instruction will cause problems. This could be worked around without too much difficulty by separating the indicators, allowing a core to be in both transaction mode (ie performing a local transaction), and recording its writes to be validated against another transaction.

Our current implementation also assumes that it is the only process running, and ignores the potential of aborts cause by external processes. While this assumption must be removed in order for transactional memory to truly be effective, this has allowed us to implement much of the basic tools and structure necessary for transactional memory.

- Assume all activity comes from a single core (will be fixed)

## 5. DESIGN CHALLENGES

While in the process of designing and implementing this feature in Palacios, we encountered a number of challenges which ultimately shaped how our implementation works. The vast majority of the difficulty arises from the need to return control to Palacios everytime a memory operation occurs. Since these normally occur extremely often with no interference, there is no existing mechanism to cause such an action. In order to do so, we chose to make use of shadow paging. By inducing page faults (namely by blindly invalidating the entire table as described above), we are able to return control to Palacios nearly every time an operation occurs.

However, additional mechanisms were still necessary: after an instruction has completed, we need to again invalidate the shadow tables in order to guarantee the next instruction will cause faults. In order to do this, we make use of a VMCall, which we insert after every instruction, as described above. In order to have the hypercall run the correct handler, the correct code must be moved into RAX prior to the execution of the call. This adds the further complication that the value of RAX must also be copied into the guest core datastructure, and restored after the VMCall has executed, in order to guarantee that the transaction code executes correctly.

In the majority of our tests, we made use of the `addq` instruction, as the decoder is able to properly determine its length and it provides us with the potential to create a variety of operations. One interesting note, is that when adding an immediate to a memory address, no detectible read is performed on the destination address: the first page fault seen after the `ifetch` is a write. Therefore, when a transaction is running, and the write is to be done on the staging page instead of main memory, the instruction will write an incorrect value on the staging page. For example, adding the immediate 1 to an address which contains a 4 will end up writing 1 ( $0 + 1$ ) to a previously blank staging page. Therefore, when adding a write to the staging page for the first time, the value in main memory must first be copied to the staging page. While the solution to this problem was not particularly complex, it stands as an example of the complexity that occurs when trying to capture all memory activity.

- Interrupts

## 6. FUTURE WORK

Future directions for this work are clear: extend the current rudimentary implementation until it is fully compatible with the Intel specifications. In particular, this means eliminating many of the limiting assumptions that were made in the presented design.

First, the limitation that instructions must fit on a page could be solved by checking the address of the instruction prior to the code injection. If the instruction is too near the edge of a page, appropriate steps could be taken to properly inject and restore the code. While this is again not overly

difficult, it does add significant complexity to the instruction injection mechanism.

The assumptions related to the staging page (that all writes fit on a page and that none of the writes have overlapping offsets), could be eliminated by no longer using the staging page as a master copy of memory for a transaction, but instead using it as a place for scratch work. When a write occurs, we pass the PTE handler the staging page, as before, but when we next exit via the hypercall, we copy the result of the write to the write set data structure. When performing a read on an address in the write set, we first copy the data to the staging page, ensuring that the read sees the newest possible version.

There are more complex behaviors described, which are not yet implemented, but so not depend on any general structural changes. For example, the Intel specification allows for nested transactions in certain situations. Implementing this would likely only require minor modifications to the state machine.

Finally, there is the issue of performance. While our implementation is able to capture and record all operations that it needs to track, it does so at the cost of up to 4 exits per instruction, which would result in a tremendous slowdown. It is likely that the necessary information could be retrieved at a less granular scale and reduce the number of exits to Palacios.

One proposed solution is to emulate all instructions within a transaction, and therefore handling the entire transaction in the host. While this would eliminate the number of exits and entries, implementing an emulator and managing the transaction are non-trivial endeavors.

## 7. CONCLUSIONS

We have implemented the preliminary components of support for the Intel model of transactional memory in the Palacios virtual machine manager. Despite a number of complex technical challenges, we have been able to create transactions which are able to log their memory activity, and either commit the data to main memory, or abort it in the case of interference. While the current implementation is subject to a number of restrictions, we have completed a number of the core components, which should greatly simplify the remainder of the task.

More generally, we have demonstrated that Palacios, and virtual machines in general, stand to provide a powerful platform for the design and testing of future hardware mechanisms. Using the power and control of Palacios, numerous other paradigms could be tested, without waiting for expensive, and possibly incomplete, implementations from hardware manufacturers. Furthermore, it could greatly simplify simulations for such hardware testing tools.

## 8. REFERENCES

- [1] C. Scott Ananian, Krste Asanovic, Bradley C. Kuszmaul, Charles E. Leiserson, and Sean Lie. Unbounded transactional memory. *IEEE Micro*, 26(1):59–69, January 2006.

- [2] Lance Hammond, Vicky Wong, Mike Chen, Brian D. Carlstrom, John D. Davis, Ben Hertzberg, Manohar K. Prabhu, Honggo Wijaya, Christos Kozyrakis, and Kunle Olukotun. Transactional memory coherence and consistency. *SIGARCH Comput. Archit. News*, 32(2):102–, March 2004.
- [3] M. Herlihy and J. Moss. Transactional memory: architectural support for lock-free data structures. *SIGARCH Comput. Archit. News*, 21(2):289–300, May 1993.
- [4] Intel Corporation. *Intel<sup>®</sup> Architecture Instruction Set Extensions Programming Reference*. Number 319433-012A. February 2012.
- [5] R. Rajwar, M. Herlihy, and K. Lai. Virtualizing transactional memory. *SIGARCH Comput. Archit. News*, 33(2):494–505, May 2005.
- [6] James Reinders. Transactional synchronization in haswell, 2012.
- [7] Nir Shavit and Dan Touitou. Software transactional memory. In *Proceedings of the fourteenth annual ACM symposium on Principles of distributed computing*, PODC '95, pages 204–213, New York, NY, USA, 1995. ACM.
- [8] Michael Solovyov and Alexander Kudryavtsev. Quix86 decoder.