

Virtualized Transactional Memory

Zachary Bischof
Northwestern University
Evanston, IL, USA
zbischof@u.northwestern.edu

Marcel Flores
Northwestern University
Evanston, IL, USA
marcel-flores@u.northwestern.edu

Maciej Swiech
Northwestern University
Evanston, IL, USA
maciejswiech2007@u.northwestern.edu

ABSTRACT

Transactional memory has long been considered as an alternative to locks for implementing parallel algorithms. In particular, it offers an optimistic approach to parallelization, and attempts to perform an operation assuming no contention will occur.

Despite the potential gains of transactional memory, it is only in the not-yet-released Haswell architecture from Intel that it has made an appearance. Palacios has presented us with a unique opportunity to emulate this Intel implementation in software.

In this paper, we present a basic implementation of transactional memory, which emulates the Haswell instructions, successfully performing memory transactions.

1. INTRODUCTION

Transactional Memory is pretty cool stuff. We decided to implement it.

2. BACKGROUND

Transactional memory has long been proposed as a solution to the programmatic complexity that arises from attempting to use mutual exclusion locks in real world implementations of parallel algorithms. We will now consider a basic overview of the general idea of transactional memory. We will then go on to describe the version that Intel has included in the specifications for the new Haswell processors.

2.1 Transactions

We define a transaction to be a sequence of memory operations that are performed by a processes such that, to all other processes, the operations seem to have been performed atomically. Furthermore, the operations must appear to the process performing the transaction as having

happened one after another, with no interference from other processes. When a transaction is finished, it either commits the results of its operations to main memory, or it aborts the transaction and throws away its changes, depending on the results of a memory validation.

In order to protect the appearance that the operations have executed serially without interference, the validation must detect if other processes have written to memory that has been used by the transaction. Likewise, in order to preserve the atomicity, the transaction must not allow other processes to interfere with the commit. Therefore, the validation must detect if other processes have written to memory that is read by the transaction (and thereby interfering with the appearance of serial execution), or other processes read memory written by the transaction (potentially interfering with atomicity). Upon an abort, a transaction may decide to attempt to run again, or revert to traditional mutex locking mechanisms. [1]

In addition to memory conflicts with other user processes, transactions may also be aborted in the event of interrupts, context switches, or other forced changes in the control flow, as they would violate the requirement that the transaction happen entirely serially.

Transactions of this form can be thought of as an optimistic approach to concurrency, contrasting with the more pessimistic approach of locks. Rather than explicitly preventing other processes from accessing memory, the transaction assumes that it will probably be able to complete unencumbered. If a conflict does occur, the transaction is able to deal with it accordingly. Locks, on the other hand, assume that conflicts are likely to occur, and directly block other processes from accessing the relevant portions of memory. Transactions therefor stand to offer a performance gain, as forced serialization can be avoided in many cases.

In addition, transactional memory stands to decrease the complexity of multi-process applications. Programmers would no longer be weighed down by keeping track of which process is holding which locks, and no longer needs to worry about deadlocks, starvation, and other side effects that can result from incorrect use of locks.

2.2 Intel Implementation

The starting with the Haswell, Intel has included the Intel Transactional Synchronization Extension, which is their implementation of transactional memory. The extension includes two interfaces: the first is Hardware Lock Elision, a system designed to work on legacy processors, and of no further interest to us. The second is Restricted Transactional Memory (RTM). This includes instructions for beginning a transactional region (and specifying the code to run in the event of a failure), aborting a transaction, and ending a transactional region. (TODO: Cite the intel manual somewhere near)

The Intel implementation generally follows the design of [1], with a few exceptions. Most notably, each transaction is only attempted a single time: in the event of a failure or abort, the code immediately resorts to the failsafe. Additionally, the transactions are checked at a granularity of cache lines, rather than actual memory address. In order to validate each transaction at the time of completion, this implementation stores a read set, the locations of memory reads from the transaction, and the write set. An abort is triggered if another processor reads from an address that appears in the write set or writes to an address in either set.

Since this is a hardware implementation, transactions may also abort if the read and write sets exceed a certain, hardware dependent, capacity. The transaction will also be aborted for any number of instructions that may interfere with control flow: such as the writing of control registers, any variety of interrupt, processor state changes, TLB control, and so on. The specification also suggests that aborts may occur in the case of self modifying code, or other unpredictable scenarios.

The Intel implementation also offers support for nested paging, however we save discussion of this topic for later work. (Maybe just don't mention this).

The Palacios virtual machine manager offers us a unique opportunity to implement transactional memory which emulates both the original design of transactional memory, as well as a number of the subtleties observed in Intel's implementation.

3. PALACIOS IMPLEMENTATION

In order to implement transactional memory in Palacios, we attempt to follow Intel's Haswell specifications, providing us with a consistent interface with which to use transactions, as well as making future comparisons to direct hardware implementations easier. We will first consider the high level implementation, in particular the state machine which governs the functioning of the transaction. We will then consider the implementation in palacios, and discuss a number of the difficulties this implementation presented.

NOTE: I'm not sure this is the best approach to describing this. Many of the decisions made in the design of the machine exist because of the complexity of trying to get back into the VMM from inside the VM, etc.

3.1 State Machine

In order to keep track of all memory reads and writes during a transaction, and to respond accordingly, we use the following state machine:

1. Begin Transaction
2. ifetch
3. Exec
4. Abort

- Diagram of the state machine?
- Note what we keep track of
- Undefined exception handler
- Mechanism by which we capture reads and writes
- General shadow page manipulations
- Staging page
- Actual trans mem handlers
- Keeping track of other cores
- Various aborts

4. DESIGN CHALLENGES

- Shadow paging subtleties
- Code injection
- Register clobbering

4.1 Current Assumptions

- AMD
- Decoder limitations
- Assume all activity comes from a single core (will be fixed)
- All memory activity from a transaction fits on a page (with no offset collisions)
- All instructions fit on a page, and overwriting the next instruction can be done on the same page.
- All interrupts (aside from the obvious ones) are ignorable

5. FUTURE WORK

- Fix major assumptions, in particular those related to pages and transaction size limitations.

6. CONCLUSIONS

- The virtual machine is a viable place to implement paradigms that do not yet exist in hardware
- While by no means fast, in our very limited transaction test space, performance was tolerable.

7. REFERENCES

- [1] M. Herlihy and J. Moss. Transactional memory: architectural support for lock-free data structures. *SIGARCH Comput. Archit. News*, 21(2):289–300, May 1993.